

101 Questions to Ask When Considering a NoSQL Database

Wednesday, June 15, 2011 at 8:08AM

Todd Hoff in nosql

You need answers, I know, but all I have here are some questions to consider when thinking about which database to use. These are taken from my [webinar What Should I Do? Choosing SQL, NoSQL or Both for Scalable Web Applications](#). It's a companion article to [What The Heck Are You Actually Using NoSQL For?](#)



Actually, I don't even know if there are a 101 questions, but there are a lot/way too many. You might want to use these questions as kind of a NoSQL [I Ching](#), guiding your way through the immense possibility space of options that are in front of you. Nothing is fated, all is interpreted, but it might just trigger a new insight or two along the way.

Where are you starting from?

A can do anything green field application?

In the middle of a project and worried about hitting bottlenecks?

Worried about hitting the scaling wall once you deploy?

Adding a separate loosely coupled service to an existing system?

What are your resources? expertise? budget?

What are your pain points? What's so important that if it fails you will fail? What forces are pushing you?

What are your priorities? Prioritize them. What is really important to you, what must get done?

What are your risks? Prioritize them. Is the risk of being unavailable more important than being inconsistent?

What are you trying to accomplish?

What are you trying to accomplish?

What's the delivery schedule?

Do the research to be specific, like Facebook did with their [messaging system](#):

Facebook chose HBase because they monitored their usage and figured out what was needed: a system that could handle two types of data patterns.

Things to Consider...Your Problem

Do you need to build a custom system?

Access patterns: 1) A short set of temporal data that tends to be volatile 2) An ever-growing set of data that rarely gets accessed 3)

High write loads 4) High throughput, 5) Sequential, 6) Random

Requires scalability?

Is availability more important than consistency, or is it latency, transactions, durability, performance, or ease of use?

Cloud or colo? Hosted services? Resources like disk space?

Can you find people who know the stack?

Tired of the data transformation (ORM) treadmill?

Store data that can be accessed quickly and is used often?

Would like a high level interface like PaaS?

Things to Consider...Money

Cost? With money you have different options than if you don't. You

can probably make the technologies you know best scale.

Inexpensive scaling?

Lower operations cost?

No sysadmins?

Type of license?

Support costs?

Things to Consider...Programming

Flexible datatypes and schemas?

Support for which language bindings?

Web support: JSON, REST, HTTP, JSON-RPC

Built-in stored procedure support? Javascript?

Platform support: mobile, workstation, cloud

Transaction support: key-value, distributed, ACID, BASE, eventual consistency, multi-object ACID transactions.

Datatype support: graph, key-value, row, column, JSON, document, references, relationships, advanced data structures, large BLOBs.

Prefer the simplicity of transaction model where you can just update and be done with it? In-memory makes it fast enough and big systems can fit on just a few nodes.

Things to Consider...Performance

Performance metrics: IOPS/sec, reads, writes, streaming?

Support for your access pattern: random read/write; sequential read/write; large or small or whatever chunk size you use.

Are you storing frequently updated bits of data?

High Concurrency vs High Performance?

Problems that limit the type of work load you care about?

Peak QPS on highly-concurrent workloads?

Test your specific scenarios?

Things to Consider...Features

Spooky scalability at a distance: support across multiple data-centers?

Ease of installation, configuration, operations, development, deployment, support, manage, upgrade, etc.

Data Integrity: In DDL, Stored Procedure, or App

Persistence design: Memtable/SSTable; Append-only B-tree; B-tree;

On-disk linked lists; In-memory replicated; In-memory snapshots;

In-memory only; Hash; Pluggable.

Schema support: none, rigid, optional, mixed

Storage model: embedded, client/server, distributed, in-memory

Support for search, secondary indexes, range queries, ad-hoc queries, MapReduce?

Hitless upgrades?

Things to Consider...More Features

Tunability of consistency models?

Tools availability and product maturity?

Expand rapidly? Develop rapidly? Change rapidly?

Durability? On power failure?

Bulk import? Export?

Hitless upgrades?

Materialized views for rollups of attributes?

Built-in web server support?

Authentication, authorization, validation?

Continuous write-behind for system sync?

What is the story for availability, data-loss prevention, backup and restore?

Automatic load balancing, partitioning, and repartitioning?

Live addition and removal of machines?

Things to Consider...The Vendor

Viability of the company?

Future direction?

Community and support list quality?

Support responsiveness?

How do they handle disasters?

Quality and quantity of partnerships developed?

Customer support: enterprise-level SLA, paid support, none

Article originally appeared on High Scalability (<http://highscalability.com/>).

See website for complete article licensing information.