


# Cell Architectures

Wednesday, May 9, 2012 at 9:15AM

Todd Hoff in Strategy

A consequence of Service Oriented Architectures is the burning need to provide services at scale. The architecture that has evolved to satisfy these requirements is a little known technique called the Cell Architecture.



A Cell Architecture is based on the idea that massive scale requires parallelization and parallelization requires components be isolated from each other. These islands of isolation are called cells. A cell is a self-contained installation that can satisfy all the operations for a [shard](#). A shard is a subset of a much larger dataset, typically a range of users, for example.

Cell Architectures have several advantages:

- Cells provide a unit of parallelization that can be adjusted to any size as the user base grows.

- Cells are added in an incremental fashion as more capacity is required.

- Cells isolate failures. One cell failure does not impact other cells.

- Cells provide isolation as the storage and application horsepower to process requests is independent of other cells.

- Cells enable nice capabilities like the ability to test upgrades, implement rolling upgrades, and test different versions of software.

Cells can fail, be upgraded, and distributed across datacenters independent of other cells.

A number of startups make use of Cell Architectures:

**Tumblr:** Users are mapped into cells and many cells exist per data center. Each cell has an HBase cluster, service cluster, and Redis caching cluster. Users are homed to a cell and all cells consume all posts via firehose updates. Background tasks consume from the firehose to populate tables and process requests. Each cell stores a single copy of all posts.

**Flickr:** Uses a federated approach where all a user's data is stored on a shard which is a cluster of different services.

**Facebook:** The Messages service has as the basic building block of their system a cluster of machines and services called a cell. A cell consists of ZooKeeper controllers, an application server cluster, and a metadata store.

**Salesforce:** Salesforce is architected in terms of pods. Pods are self-contained sets of functionality consisting of 50 nodes, Oracle RAC servers, and Java application servers. Each pod supports many thousands of customers. If a pod fails only the users on that pod are impacted.

The key to the cell is you are creating a scalable and robust MTBF friendly service. A service than can be used as a bedrock component in a system of other services coordinated by a programmable orchestration layer. It works just as well in a data center as in a cloud. If you are looking for a higher level organization pattern, the Cell Architecture is a solid choice.

## Related Articles

## Startups Are Creating A New System Of The World For IT

---

Article originally appeared on High Scalability (<http://highscalability.com/>).

See website for complete article licensing information.